

Workflow of Supervised Techniques

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**,

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**, possibly matched in some way to nature of problem.

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**, possibly matched in some way to nature of problem.
- 3 Adjust algorithm for **optimal performance**,

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**, possibly matched in some way to nature of problem.
- 3 Adjust algorithm for **optimal performance**, perhaps using **validation set** and/or some kind of **cross-validation**.

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**, possibly matched in some way to nature of problem.
- 3 Adjust algorithm for **optimal performance**, perhaps using **validation set** and/or some kind of **cross-validation**.
- 4 Report accuracy in **test set**,

Workflow of Supervised Techniques

- 1 Obtain/code the **training set** and decide on relevant features (preprocess).
- 2 Decide on the **algorithm**, possibly matched in some way to nature of problem.
- 3 Adjust algorithm for **optimal performance**, perhaps using **validation set** and/or some kind of **cross-validation**.
- 4 Report accuracy in **test set**, possibly combine with other learners in **ensemble**.

Notes and Issues

Notes and Issues

Supervised techniques are about **learning** relationship between X and labeled data.

Notes and Issues

Supervised techniques are about *learning* relationship between X and labeled data. Often used interchangeably with *machine learning*: idea that computers could 'learn' relationships without specific programming.

Notes and Issues

Supervised techniques are about *learning* relationship between X and labeled data. Often used interchangeably with *machine learning*: idea that computers could 'learn' relationships without specific programming.

Typically used when $p \gg n$:

Notes and Issues

Supervised techniques are about *learning* relationship between X and labeled data. Often used interchangeably with *machine learning*: idea that computers could 'learn' relationships without specific programming.

Typically used when $p \gg n$: number of parameters (e.g. coefficients on terms) is far larger than number of observations (e.g. speakers or documents).

Notes and Issues

Supervised techniques are about **learning** relationship between X and labeled data. Often used interchangeably with **machine learning**: idea that computers could 'learn' relationships without specific programming.

Typically used when $p \gg n$: number of parameters (e.g. coefficients on terms) is far larger than number of observations (e.g. speakers or documents).

- results in general **curse of dimensionality** wherein feature matrix is large (e.g. 100k columns) and **sparse** and thus obtaining meaningful estimates is difficult.

Notes and Issues

Supervised techniques are about **learning** relationship between X and labeled data. Often used interchangeably with **machine learning**: idea that computers could 'learn' relationships without specific programming.

Typically used when $p \gg n$: number of parameters (e.g. coefficients on terms) is far larger than number of observations (e.g. speakers or documents).

→ results in general **curse of dimensionality** wherein feature matrix is large (e.g. 100k columns) and **sparse** and thus obtaining meaningful estimates is difficult.

So techniques may require careful tuning of **regularization parameters** to obtain good performance.

Notes and Issues II

Notes and Issues II

Once we have the **training set** we face a dilemma...

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly),

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements.

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set,

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set,

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance.

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**,

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**, but may induce **bias** in the sense that we miss important relationships in the data.

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**, but may induce **bias** in the sense that we miss important relationships in the data.
→ we have **underfit** to our training set,

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
- 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**, but may induce **bias** in the sense that we miss important relationships in the data.
→ we have **underfit** to our training set, and may do poorly on our test set.

Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
 - 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**, but may induce **bias** in the sense that we miss important relationships in the data.
→ we have **underfit** to our training set, and may do poorly on our test set.
- So managing the ***bias-variance tradeoff*** is a key element of supervised learning,


Notes and Issues II

Once we have the **training set** we face a dilemma...

- 1 we can use our technique to fit a very complicated model to this data (perfectly), including noisy elements. This avoids **bias**, but it incurs **variance** (when we move to the test set).
→ we have **overfit** to our training set, and may do poorly on our test set.
 - 2 we can use be more relaxed about the fit of our algorithm in the training set, and accept some imperfections in performance. This avoids high **variance**, but may induce **bias** in the sense that we miss important relationships in the data.
→ we have **underfit** to our training set, and may do poorly on our test set.
- So managing the **bias-variance tradeoff** is a key element of supervised learning, and we may need to **tune** our algorithms with that in mind.

Bias-Variance Tradeoff (Hastie et al, p38)

Bias-Variance Tradeoff (Hastie et al, p38)



Cross-Validation

Cross-Validation

Want to properly estimate model performance (bias and especially variance).

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- 1 divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'),

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- 1 divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- 1 divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- 2 repeat the following k times (folds):

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**
 - ③ test on the validation set,

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**
 - ③ test on the validation set, record **prediction error**

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**
 - ③ test on the validation set, record **prediction error**
- ③ Average over runs to get prediction error estimate.

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**
 - ③ test on the validation set, record **prediction error**
- ③ Average over runs to get prediction error estimate.

→ Can follow same steps for models of different specifications; variant on this approach can be used for **model selection**, directly.

Cross-Validation

Want to properly **estimate model performance** (bias and especially variance).

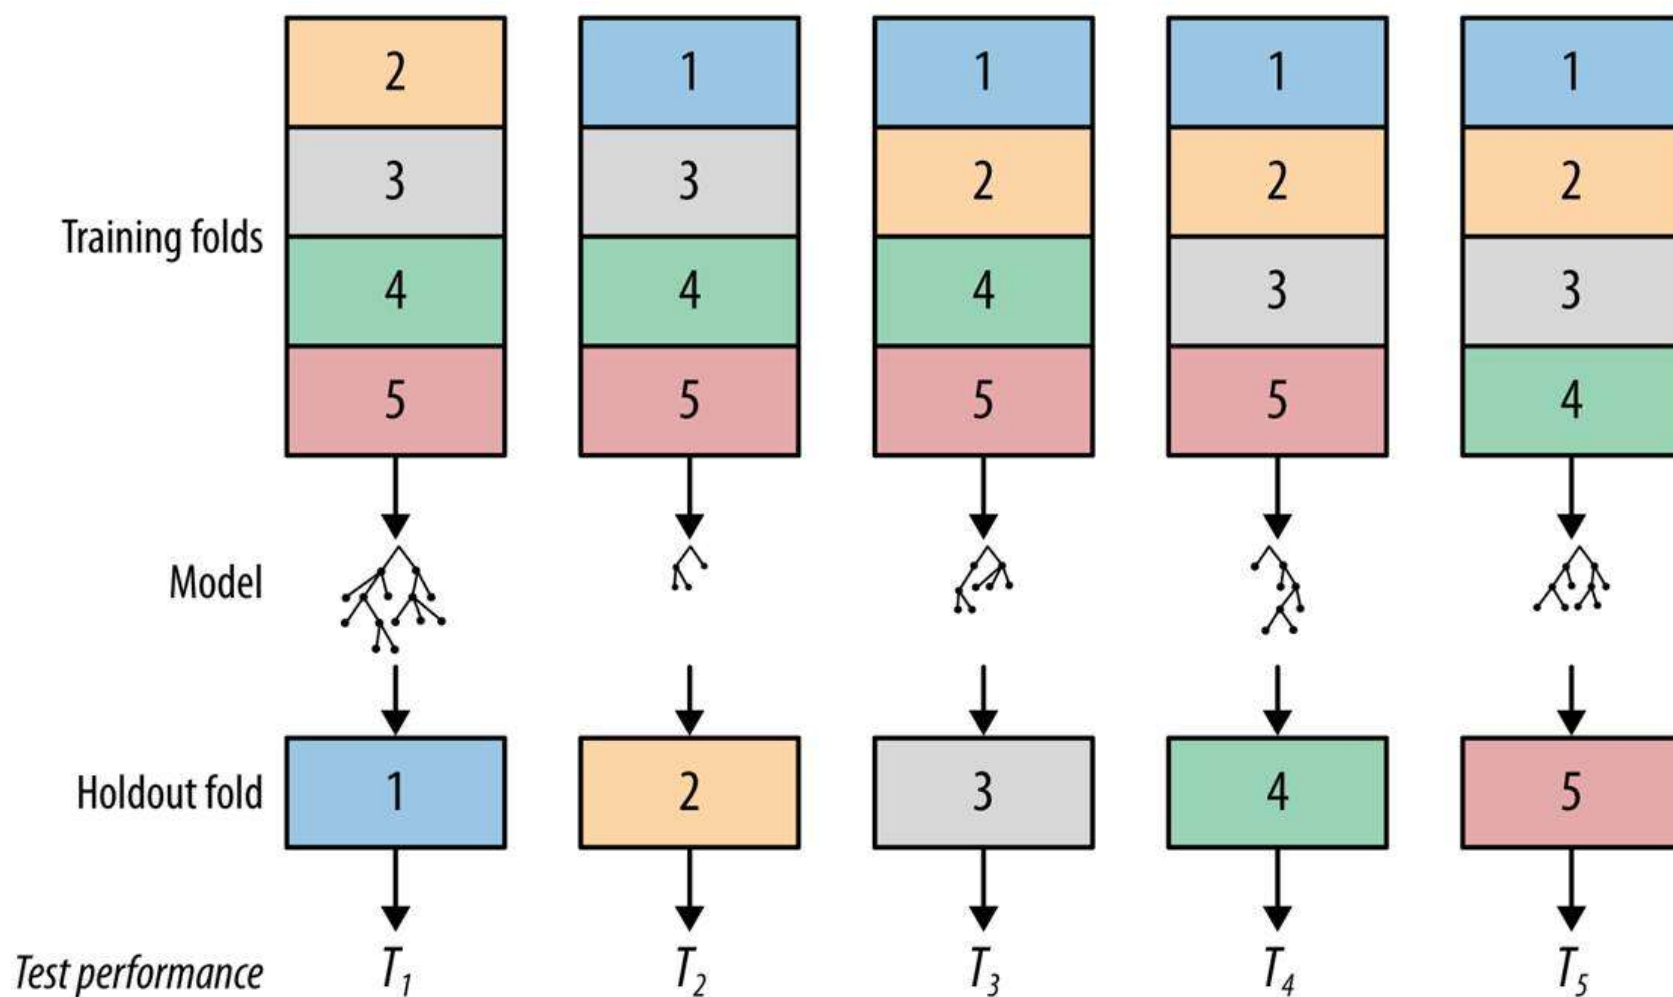
Popular and efficient approach is **k -fold cross validation**, esp when we don't have enough data to form a completely separate **validation** set.

- ① divide all data into k equal size chunks ($k = 10$ is common; $k = n$ is 'leave one out'), and set the parameter(s) of model at particular value,
- ② repeat the following k times (folds):
 - ① grab one of the k chunks as a **validation set** (each only used once)
 - ② grab the other $k - 1$ chunks as a **training set**
 - ③ test on the validation set, record **prediction error**
- ③ Average over runs to get prediction error estimate.

→ Can follow same steps for models of different specifications; variant on this approach can be used for **model selection**, directly.

Graphically

Graphically



Mean and standard deviation of test sample performance